

Implement a Key-Value Store

Using AWS S3

DISTRIBUTED COMPUTING SYSTEMS (X_400130)
LARGE LAB EXERCISE

S. Talluri

V1.0

1 Learning Objectives

1. Design and conduct experiments to understand the behavior of real-world systems.
2. Implement new systems to take advantage of real-world systems to provide new features.
3. Analyze experiment data to gain new information about the system you implemented.

This exercise combines elements of Chapter 4 (communication), Chapter 6 (synchronization), Chapter 7 (consistency and replication), and Chapter 8 (fault tolerance) of the course text book (Andrew S. Tanenbaum and Maarten van Steen, Distributed Systems: Principles and Paradigms, third edition, Prentice Hall, 2017).

2 Assignment Description

Key-Value Stores are an important component of many distributed systems. Systems like databases, games, web applications, and schedulers store their persistent state in a key value store so that multiple components of the system see the same data, and the data is stored in a fault tolerant manner.

KV stores achieve fault tolerance by usually persisting the data to local disk. But, cloud virtual machines such as those from AWS EC2 do not come with local disk. They usually use remote byte addressable storage such as AWS EBS. With EBS, users pay for the time they use the allocated disk. This is different from cloud native storage services such as S3 for which users pay per read or write.

In this project, you will understand the overhead imposed by services such as S3 for small writes. You will implement a KV layer on top of S3 to mitigate this overhead. You will also learn about log structuring to store data on media which disallows random writes.

Related work: Check "BlueSky: A Cloud-Backed File System for the Enterprise".

2.1 Requirements

Mandatory Requirements

- ☐ Setup a benchmark to measure the performance of uploading small key-value pairs (32B - 1KB) to S3.
- ☐ Implement a KV store that uses AWS S3 as a persistent log to persist data.

- ☐ Design and implement experiments to demonstrate the performance and fault tolerance characteristics of your system.

Bonus Features – You Will Receive Bonus Points for at Most 4 Implemented Features

- ☐ Implement garbage collection to remove old versions of KV pairs.
- ☐ Implement compaction support scans of sequential keys.
- ☐ Your **own** idea: we support students pursuing their own interesting idea. In fact, we encourage teams to think of additional features they can add to a product showing the conceptual and engineering skills. If you think of an interesting feature, please discuss it with a TA first.

2.2 Deadlines

1. (mandatory) November 5, 2021 (week 44): Form a group and pick a lab project. Each project has a group created on Canvas, make sure all members enroll in that group.
2. (recommended) November 12, 2021 (week 45): Discuss and finalize your system's functional and non-functional requirements. It is recommended to discuss this with your supervisor.
3. (recommended) November 19, (week 46): Discuss with the supervisor your proposed design and plan for implementation.
4. (mandatory) Anytime before December 9, 2021 (week 49): Demonstrate to the supervisor the system implemented for the lab exercise.
5. (recommended) December 15, 2021 (week 50): Demo Day for the selected groups.
6. (mandatory) December 21, 2021 (week 51): Turn in the lab-exercise report on canvas.

To achieve the above-mentioned deadlines, your assignment can be divided into three parts:

Weeks 2–3:

- Analyze the system functionality and summarize its requirements.
- Analyze the (potential) system implementation and summarize your design's key features. Discuss with your team-mate(s) the alternatives for the design choices you have made, and summarize them in your report. Try to have the resulting requirements and design document (maximum 2 pages) approved by the supervisor before you start the actual implementation.
- Implement a part of the key features

Weeks 4–5:

- Finalize the design.
- Implement all key features.
- Test your system's implementation for different features: correctness, consistence, scalability etc.
- Per feature, consider how you are going to test its correctness and report on it.

Weeks 6–7:

- Implement and execute the test plans.

- Demonstrate your **working** prototype/product to the supervisor. Any feature that you claim it has, must be demonstrated. We recommend teams to make scripts that demonstrate each feature, to avoid having to setup everything during the demo. The demo can take at most 30 minutes, use this time wisely!
- Finish your report by including the achieved results and your conclusions, and send it to the supervisor.

Note: the estimated time required for the completion of this assignment (the large exercise) is 80 hours per team member. We expect the work to be equally divided among all team members.

Notes:

- Try not to exceed, for any experiment, 10 hours of work.
- You can leverage the same development and setup for several experiments. In this case, report the time spent for the shared part only for the first experiment and explain the large amount of time spent for it in the report.
- You can look at related scientific material to see how they tested their system, using something similar usually gives you a good basis.

3 Report Structure

The report should have the following structure:

1. *Report title, authors, and support cast* (lab supervisor and course instructors). For each person in your team, give name and contact information (email).
2. *Abstract*: a description of the problem, system description, analysis overview, and one main result. Size: one paragraph with at most 150 words.
3. *Introduction* (recommended size, including points 1 and 2: 1 page): describe the problem, the existing systems and/or tools about which you know (related work), the system you are about to implement, and the structure of the remainder of the article. Use one short paragraph for each.
4. *Background* on the application (recommended size: 0.5 pages): describe your product or the system you are working with (1 paragraph) and its requirements (1 paragraph per each of consistency, scalability, fault-tolerance, performance, etc.).
5. *System Design* (recommended size: 1.5 pages)
 - System overview: describe the design of your system or research, including the system operation, fault tolerance, and scalability components (which correspond to the homonym features required in this project).
 - (Optional, for bonus points, see Section 4) Additional (system) features: describe each additional feature, one sub-section per feature.
6. *Experimental Results* (recommended size: 1.5 pages)
 - (a) *Experimental setup*: describe the working environments (DAS, Amazon EC2, etc.), the general workload and monitoring tools and libraries, other tools and libraries you have used to implement and deploy your system, and tools and libraries used to conduct your experiments.

- (b) *Experiments*: describe the experiments you have conducted to analyze each (system) feature, such as consistency, scalability, fault-tolerance, and performance. Analyze the results obtained for each system feature. Use one sub-section per experiment (or feature). In the analysis, also report:
 - i. Service metrics of the experiment, such as runtime and response time of the service, etc. and describe what it measures.
 - ii. (optional) Usage metrics and costs of the experiment where applicable.
7. *Discussion* (recommended size: 1 page): summarize the main findings of your work and discuss the tradeoffs inherent in the design of your system. Should you use a distributed system to implement the requested system? Try to extrapolate from the results reported in Section 6b for system workloads that are orders of magnitude higher than what you have tried in real-world experiments.
8. *Conclusion*
9. Appendix A: Time sheets (see Section 3.1)

3.1 Document the Time You Spend

You should report on the time it takes to conduct each major part of the assignment. Specifically, report:

- the `total-time` = total amount of time spent in completing the assignment (the large exercise).
- the `think-time` = total amount of time spent in thinking about how to solve the assignment (the large exercise).
- the `dev-time` = total amount of time spent in developing the code needed to solve the assignment (the large exercise).
- the `xp-time` = total amount of time spent in experiments for the assignment (the large exercise).
- the `analysis-time` = total amount of time spent in analyzing the results of the experiments for the assignment (the large exercise).
- the `write-time` = total amount of time spent in writing this report
- the `wasted-time` = total amount of time spent in activities related to the assignment (the large exercise), but which cannot be charged as think-time, dev-time, xp-time, analysis-time, or write-time.

4 Scoring

1. Quality of the report [9, 10] (presentation, style of writing, graphing, requirements analysis, design, analysis of results): +2,000 points.
2. Technical quality of the systems work (basic system features are supported): +2,000 points.
3. Bonuses
 - Making your source code open-source and your report public: +100 points.
 - Additional system features (feature and report): usually +125 points per feature (see Section 2.1).

- Excellent report (writing, graphing, description of system): at most +250 points.
- Excellent analysis (design of experiments, figures, analysis of results): at most +250 points.

4. Limits on bonuses

- The total additional points for this exercise can amount to at most +1,000 points, representing additional features, excellent report and/or analysis bonuses, etc.
- The additional features (see section 2.1) can amount to at most +500 points. Pay attention that some exercises limit the amount of bonus features that you can get credits for. Usually the amount of bonus points obtainable is equal per feature, but this may differ between assignments.

Note: writing a good technical report is scored equally to writing code for a system that works! This means that you should schedule time to write a good report, and for one revision to respond to our feedback.

5 Use of Existing Technology

You are not limited for this exercise to any technology, but:

- You CAN use the machines provided by our system DAS-4 or DAS-5, or by Amazon EC2 or another IaaS cloud (e.g. Azure, Google Cloud). For Amazon EC2, you can use the free resources provided through the Free Usage Tier [<http://aws.amazon.com/free/>]; if you use the m1.small (paid) instances of Amazon EC2, the estimated cost for this exercise does not exceed 8 EUR¹.
- You CAN use libraries that already provide a mandatory functionality required in Section C, as long as you implement a bonus feature listed in Section C in exchange AND after approval of the supervisor. If you only implement the mandatory functionalities, you CAN NOT use such libraries.

Examples of libraries that are fine to use when implementing only the mandatory functionalities are:

- Google RPC,
- Python's default communication libraries (sockets),
- Java Akka.

Examples of libraries NOT to use UNLESS you implement a bonus feature and you received approval of the supervisor are:

- Hazelcast, it provides fault-tolerance automatically,
- Frameworks such as Hadoop/Spark that automate fault-tolerance and load-balancing.

When using an external library, give due credit in your code and report (see also Section 6).

When in doubt about technology you intend to use, consult with the supervisor.

¹At 8 Euro-cents per hour, assuming a total workload of 100 hours. Mileage may vary.

6 Anti-Fraud Policy (Zero-Tolerance)

Our anti-fraud policy is simple: zero-tolerance, within the limits set by VU Amsterdam. We will pursue each case of potential fraud, and will use to the maximum extent the means provided by VU Amsterdam to prevent, then to discover and punish (attempts to) fraud. You can learn more about how to prevent that you commit fraud via a discussion with the “studieadviseur”, from Appendix 2 in the Regulations and Guidelines regarding examinations FEWEB [1], or even from international sources such as Harvard’s guidelines on avoiding plagiarism [2].

References

- [1] AMSTERDAM, V. Student charter 2017-2018, 2018.
- [2] HARVARD. Avoiding plagiarism, 2019.